

Lecture 7: Minimum Spanning Trees & Shortest Paths

26 October 2019

Lecturer: Chaya Hiruncharoenvate

Scribe: -

* Recap Greedy Algorithm from last week

Today we will discuss another important graph problem called Minimum Spanning Tree or MST. MST has many applications in real-world such as networking, transportation, etc.

1 Minimum Spanning Trees

Definition 1.1. Let $H = (V, T)$ be a subgraph of an undirected graph $G = (V, E)$. H is a **spanning tree** of G if H is both acyclic and connected.

Spanning trees have many interesting properties. The following statements are equivalent:

- H is a spanning tree of G
- H is acyclic and connected
- H is connected and has $n - 1$ edges
- H is acyclic and has $n - 1$ edges
- H is minimally connected: removal of any edge disconnects it
- H is maximally acyclic: addition of any edge creates a cycle
- H has a unique simple path between every pair of nodes

Now that we have defined what a spanning tree is. Let's talk about MST.

Definition 1.2. Given a connected, undirected graph $G = (V, E)$ with edge costs c_e , a **minimum spanning tree** (V, T) is a spanning tree of G such that the sum of the edge costs in T is minimized.

[Draw a simple connected graph and identify an MST]

Instructor note: Point out that MST is not always unique.

1.1 Greedy Approach

Consider the following greedy approach: start with an empty graph and repeatedly add the next lightest edge that doesn't produce a cycle. In other words, it constructs the tree edge by edge and simply picks whichever edge is cheapest at the moment while avoiding making a cycle.

This algorithm is known as **Kruskal's algorithm**. How do we know that Kruskal's tree is optimal? We will consider something called the **cut property**.

Definition 1.3. A *cut* S is a partition of nodes into two nonempty subsets: S and $V - S$

Definition 1.4. The *cutset* of a cut S is a set of edges with exactly one endpoint in S .

[Draw a graph, define a cut S and identify the cutset of S]

Proposition 1.5. *A cycle and a cutset intersect in an even number of edges*

Proof: Consider a cycle and any cut S . Idea: when there is a edge of the cycle going into S there must another edge going out of S . ■

Lemma 1.6 (Cut property). *Suppose edges X are part of a minimum spanning tree of $G = (V, E)$. Pick a cut S for which no edges of X in the cutset of S , and let e be the lightest edge across this partition. Then $X \cup \{e\}$ is part of some MST.*

Proof: Suppose edges X are part of an MST T . If e is a part of T , then we're done. If not, we will show that there exists another MST T' that contains $X \cup \{e\}$. Consider $T \cup \{e\}$. Since T is connected, there is already a path between the two endpoints of e . So, when e is added, there will be a cycle. Along this cycle, there must be another edge e' across the cut $(S, V - S)$. If we remove this e' , we are left with $T' = T \cup \{e\} - \{e'\}$. We know that T' is a tree because T' is connected and has $n - 1$ edges. Moreover, we know that T' is a minimum spanning tree because:

$$\text{weight}(T') = \text{weight}(T) + c_e - c_{e'}$$

Since both e and e' are in the cutset of S , we know that $c_e \leq c_{e'}$. So, $\text{weight}(T') \leq \text{weight}(T)$. Since T is a minimum spanning tree, it must be the case that $\text{weight}(T') = \text{weight}(T)$. So, T' is also a MST. ■

In summary, the cut property says is that it is always safe to add the lightest edge across any cut (that is, between a vertex in S and one in $V - S$), provided the current partial tree X has no edges across the cut.

Corollary 1.7. *Kruskal's algorithm is optimal*

Proof: At any given moment, Kruskal's algorithm maintains a partial solution X , a collection of connected, acyclic components. The next edge e to be added connects two of these components; call them T_1 and T_2 . Since e is the lightest edge that doesn't produce a cycle, it is certain to be the lightest edge between T_1 and $V \setminus T_1$ and therefore satisfies the cut property. So, $X \cup \{e\}$ is a part of some MST. ■

Implementation: $O((m + n) \log m)$ with the help of disjoint sets that allow us to do cycle checking very fast (amortized cost closes to $O(1)$).

1.2 Prim's algorithm

If we think about it, cut property tells us that any algorithm of the following form should work:

```
X = { } // edges picked so far
repeat until |X| = |V| - 1:
    pick a cut S for which X has no edges in the cutset of S
    let e in E be the minimum-weight edge between S and V \ S
    X = X + {e}
```

An alternative to Kruskal's algorithm called Prim's algorithm. Here's how Prim's algorithm works. Each step the algorithm maintains a subtree X and choose S to be the vertices of X .

Prim(G)

- Choose an initial node u
- $X = \{\}$
- $cost(v) = \infty$ for all v
- $prev(v) = NULL$ for all v
- $cost(u) = 0$
- $Q = PriorityQueue(V, cost)$
- While Q is not empty,
 - $v = deleteMin(Q)$
 - For each $(v, z) \in E$:
 - * update $cost(z) = \min(cost(z), cost(v) + w(v, z))$
 - * update $prev(z)$ as needed
 - * Decrease key z in Q as needed
 - $X = X + (prev(v), v)$ (except for the u)
- return X

Running time: $O((m + n) \log n)$

2 Single-Source Shortest Paths

Given a weighted, directed graph $G = (V, E)$, with weight function $w : E \rightarrow \mathbb{R}$ mapping edges to real-valued weights.

Definition 2.1. The weight $w(p)$ of path $p = \langle v_0, v_1, \dots, v_k \rangle$ is the sum of the weights of its constituent edges:

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$$

Definition 2.2. The shortest-path weight $\delta(u, v)$ from u to v is:

$$\delta(u, v) = \begin{cases} \min\{w(p) : u \rightarrow^p v\} & \text{if there is a path from } u \text{ to } v \\ \infty & \text{Otherwise} \end{cases}$$

Definition 2.3. A **shortest path** from u to v is any path p such that $w(p) = \delta(u, v)$

In this lecture, we will focus on one of the many shortest-paths problem known as the **single-source shortest paths** problem. Given a graph $G = (V, E)$, we want to find a shortest path from a given source node $s \in V$ to each node $v \in V$.

If you think about it, an algorithm that solves SSSP will be able to solve the following variants as well:

- Single-destination shortest-paths problem
- Single-pair shortest-path problem
- All-pairs shortest-paths problem

3 Dijkstra's Algorithm

Dijkstra's algorithm solves the single-source shortest-paths problem on a weighted, directed graph $G = (V, E)$ for the case in which all edge weights are nonnegative. Therefore, we assume that $w(u, v) \geq 0$ for each edge $(u, v) \in E$.

Dijkstra's algorithm maintains a set S of vertices whose final shortest-path weights from the source s have already been determined. The algorithm repeatedly selects the vertex $u \in V - S$ with the **minimum** shortest-path estimate, adds u to S , and relaxes all edges leaving u .

In the following implementation, we use a min-priority queue Q of vertices, keyed by their d values.

DIJKSTRA(G, w, s):

1. Initialize-Single-Source(G, s)
2. $S = \{\}$
3. $Q = V$
4. While Q is not empty
 - $u = DELETEMIN(Q)$
 - $S = S + \{u\}$
 - For each edge $(u, v) \in E$: RELAX(u, v, w)

INITIALIZE-SINGLE-SOURCE(G, s):

1. For v in V
 - $d[v] = \infty$
 - $\pi[v] = NULL$
2. $d[s] = 0$

RELAX(u, v, w):

1. if $d[v] > d[u] + w(u, v)$ then
 - $d[v] = d[u] + w(u, v)$
 - $\pi[v] = u$


Here is example of how Dijkstra's algorithm work step-by-step:

Dijkstra's is actually using a greedy strategy to build up the solution. In each iteration, it chooses the lightest or cheapest edge to add to S . The correctness of Dijkstra's heavily relies on the nonnegative weight assumption.

Theorem 3.1. *Dijkstra's algorithm, run on a weighted, directed graph $G = (V, E)$ with nonnegative weight function w and source s , terminates with $d[u] = \delta(s, u)$ for all $u \in V$.*

Proof: We use the following loop invariant:

At the start of each iteration of the **while loop** $d[v] = \delta(s, v)$ for all $v \in S$



Initialization: Initially, S is empty and so the invariant is trivially true.

Maintenance: We wish to show that in each iteration $d[u] = \delta(s, u)$ for the node u added to S . For the sake of contradiction, let u be the first vertex $d[u] \neq \delta(s, u)$ when it is added to S . We know that $u \neq s$ because s is the first node added to S and that $d[s] = \delta(s, s) = 0$. Because $u \neq s$, when u is added to S , S must be non-empty. There must be some path from s to u cause $d[u] \neq \infty$. So there must be a shortest path p from s to u . Consider the first node y along the path p such that $y \in V - S$ and let $x \in S$ be a predecessor of y along p . Now, we can decompose p into $p_1 = \langle s, \dots, x \rangle$ and $p_2 = \langle y, \dots, u \rangle$.

We claim that $d[y] = \delta(s, y)$ when u is added to S . Since u is the first vertex that $d[u] \neq \delta(s, u)$ when added to S , we know $d[x] = \delta(s, x)$. Also, $x \in S$ and x is added to S and, the edge (x, y) is relaxed and that would set $d[y] = \delta(s, y)$.

Since $y \notin S$ and y appears before u on a shortest path p and all edges have nonnegative weights, we have $\delta(s, y) \leq \delta(s, u)$ and, thus,

$$d[y] = \delta(s, y) \leq \delta(s, u) \leq d[u]$$

Since both $u, y \notin S$ and u was chosen before y , it must be the case that

$$d[y] = \delta(s, y) = \delta(s, u) = d[u]$$

This contradicts our choice of u . Thus, we conclude that $d[u] = \delta(s, u)$ when u is added to S .

Termination: At termination, Q is empty. So $S = V$. Thus, $d[u] = \delta(s, u)$ for all $u \in V$. ■

Running time: Running time of Dijkstra's depends on the implementation of the priority queue. For example, a binary heap implementation yields $O(\log n)$ for each DELETE-MIN and DECREASE-KEY. Since DECREASE-KEY will be called at most $|E|$ times and DELETE-MIN $|V|$ times, the total time is $O((|V| + |E|) \log |V|)$ or $O(|E| \log |V|)$. The running time of $O(|V| \log |V| + |E|)$ is possible with a Fibonacci heap implementation.

Final Note: Dijkstra's is very much like Prim's in that they grow the graph by adding the next lightest edge.

4 Bellman-Ford Algorithm

We have shown that Dijkstra's works on graphs containing nonnegative edges. What if a graph contains negative edges?

Dijkstra's algorithm works in part because the shortest path from the starting point s to any node v must pass exclusively through nodes that are closer than v . This no longer holds when edge lengths can be negative.

So, what do we do when a graph contains negative edges? To answer this, let's take a particular high-level view of Dijkstra's algorithm. A crucial invariant is that $d[]$ values it maintains are always either overestimates or exactly correct. They start off at ∞ , and the only way they ever change is by updating along an edge:

```
def relax((u,v) in E):
    d[v] = min{d[v], d[u]+w(u,v)}
```

Here are a few observations regarding this update operation.

- This update can be applied repeatedly without hurting our shortest path.
- Dijkstra's can be thought of a sequence of updates in a particular fashion.
-
- $s \rightarrow u_1 \rightarrow u_2 \dots \rightarrow t$ Suppose P is a shortest path from s to t . It should not contain duplicate, hence the longest possible that path P can only have a maximum of $|V| - 1$ edges.

So, here is an idea. Let's do update on all edges for $|V| - 1$ times! So we can be sure that we cover all possibilities. This algorithm is called **Bellman-Ford** algorithm.

```
Bellman-Ford(V, E, l, s) {
    For all u in V: distance(u) = infinity
    d[s] = 0
    repeat n-1 times:
        for all e in E:
            relax(e)
}
```

The running time of Bellman-Ford is $O(|V||E|)$.

5 Negative cycles

When a graph contains a negative cycle i.e. a cycle whose total weight is negative, the shortest-paths problem doesn't make much sense (because you could loop over the cycle to get a shorter and shorter path).

How do we detect if a graph contains negative cycles? This is fairly simple. Bellman-Ford can help finding negative cycles. Instead of stopping after $|V| - 1$ rounds, perform an extra round of update. If some of $d[]$ reduces, then we know there must be a negative cycle.